

FDE Challenge Week 2: The Automaton Auditor

Interim Submission Report

Orchestrating Deep LangGraph Swarms for Autonomous Governance

Nahom Desalegn Adisu

Submission Date: February 26, 2026

Challenge Track: Week 2 — The Automaton Auditor

Repository: github.com/nahdes/automation-auditor

Abstract

This interim report documents the architectural decisions, implementation progress, and known gaps for the **Automaton Auditor** — a hierarchical LangGraph agent swarm designed to autonomously audit code repositories and architectural reports. The system implements a “Digital Courtroom” paradigm with three specialized layers: (1) *Detectives* for forensic evidence collection via AST parsing and sandboxed git operations, (2) *Judges* for dialectical evaluation using distinct Prosecutor/Defense/TechLead personas, and (3) a *Chief Justice* for deterministic synthesis of final verdicts. This report details the rationale for Pydantic-based state management, the structure of AST-based code analysis, the sandboxing strategy for safe tool execution, and provides concrete plans for completing the judicial layer and synthesis engine. A StateGraph diagram illustrates the planned parallel fan-out/fan-in execution topology.

Contents

1 Architecture Decisions	2
1.1 Why Pydantic Over Plain Python Dicts	2
1.2 AST Parsing Strategy for Code Analysis	2
1.3 Sandboxing Strategy for Safe Tool Execution	3
2 Known Gaps and Concrete Plans	4
2.1 Judicial Layer: Current Status and Plan	4
2.2 Synthesis Engine: Current Status and Plan	5
3 Planned StateGraph Flow	5
3.1 Architecture Overview	5
3.2 StateGraph Diagram	6
3.3 Key Execution Characteristics	6
4 Conclusion and Next Steps	6
4.1 Summary of Progress	6
4.2 Immediate Next Steps (This Week)	6
4.3 Risk Mitigation	6

A Appendix: Representative Code Snippets	7
A.1 AgentState Definition with Reducers	7
A.2 AST Visitor for Graph Structure Detection	7

1 Architecture Decisions

1.1 Why Pydantic Over Plain Python Dicts

Decision: All state objects, evidence containers, and output schemas are defined using Pydantic `BaseModel` with `TypedDict` for the LangGraph `AgentState`.

Rationale:

- **Runtime Validation:** Pydantic enforces type constraints at instantiation, catching errors early (e.g., `score: int = Field(ge=1, le=5)` prevents invalid scores).
- **Structured Output Enforcement:** Judge nodes use `.with_structured_output(JudicialOpinion)` to guarantee JSON-compliant responses, reducing hallucination risk.
- **Parallel-Safe State Reduction:** Using `Annotated[Dict, operator.ior]` and `Annotated[List, operator.add]` ensures concurrent detective/judge branches do not overwrite each other's data.
- **Self-Documenting Schemas:** Field descriptions serve as inline documentation and enable automatic OpenAPI/JSON Schema generation for debugging.
- **Serialization Consistency:** Pydantic's `model_dump()` and `model_validate()` provide reliable JSON/Markdown conversion for audit reports.

Trade-offs Considered:

- *Plain dicts*: Faster to write but prone to silent type errors and harder to debug in multi-agent flows.
- *dataclasses*: Good for immutability but lack runtime validation and JSON serialization utilities.
- *attrs*: Powerful but adds external dependency complexity without significant advantage over Pydantic v2.

Implementation Example:

```

1 from pydantic import BaseModel, Field
2 from typing import Optional, List
3
4 class Evidence(BaseModel):
5     goal: str = Field(description="What this evidence verifies")
6     found: bool = Field(description="Whether artifact exists")
7     content: Optional[str] = Field(default=None)
8     location: str = Field(description="File path or commit hash")
9     rationale: str = Field(description="Confidence rationale")
10    confidence: float = Field(ge=0.0, le=1.0)
11    tags: List[str] = Field(default_factory=list)

```

Listing 1: Pydantic Evidence Model (src/state.py)

1.2 AST Parsing Strategy for Code Analysis

Decision: Use Python's built-in `ast` module (not regex) to parse and analyze repository code structure.

Rationale:

- **Robustness:** AST parsing handles syntactically valid Python regardless of formatting, comments, or string literals that would break regex.
- **Semantic Understanding:** Can distinguish between `StateGraph` instantiation vs. mere string occurrence, and extract class hierarchies, function calls, and import graphs.
- **Extensibility:** The `ASTVisitor` pattern allows adding new forensic protocols (e.g., detecting `operator.add` reducers) without rewriting parsers.
- **Security:** No execution of untrusted code — pure static analysis.

Implementation Structure:

1. `ASTVisitor` class extends `ast.NodeVisitor` to collect:
 - Class definitions with base classes (for Pydantic/TypedDict detection)
 - Function calls (for `add_edge`, `with_structured_output` detection)
 - Import statements (for LangGraph/Pydantic dependency verification)
 - Assignment statements (for reducer pattern detection)
2. `parse_python_file()` wraps AST parsing with error handling for syntax errors.
3. `scan_directory_for_python()` recursively traverses repo, excluding virtual environments.
4. Forensic protocols (`analyze_state_management()`, `analyze_graph_structure()`) query the visitor's collected metadata.

Limitations:

- Cannot analyze dynamically generated code or code requiring runtime imports.
- Does not execute code — cannot verify functional correctness, only structural presence.
- Language-specific: Currently Python-only; multi-language repos would require tree-sitter or language-specific parsers.

1.3 Sandboxing Strategy for Safe Tool Execution

Decision: All git operations and file system interactions run inside `tempfile.TemporaryDirectory()` with `subprocess.run()` (never `os.system()`).

Rationale:

- **Isolation:** Cloned repositories cannot affect the auditor's working directory or host system.
- **Automatic Cleanup:** `TemporaryDirectory` ensures disk space is reclaimed even on errors.
- **Controlled Execution:** `subprocess.run(capture_output=True, timeout=N)` prevents hanging processes and captures stderr for debugging.
- **Input Sanitization:** Repository URLs are validated before use; no shell interpolation prevents injection attacks.
- **Error Handling:** All subprocess calls wrap exceptions and return structured Evidence objects with failure rationale.

Implementation Pattern:

```

1 def clone_repo_sandboxed(repo_url: str, depth: int = 50, timeout: int = 120):
2     tmpdir = tempfile.TemporaryDirectory(prefix="auditor_clone_")
3     try:
4         result = subprocess.run(
5             ["git", "clone", "--depth", str(depth), repo_url, tmpdir.name],
6             capture_output=True, text=True, timeout=timeout
7         )

```

```

8     if result.returncode != 0:
9         tmpdir.cleanup()
10        return None, None
11    return tmpdir.name, tmpdir # Caller must cleanup tmpdir
12 except Exception:
13     tmpdir.cleanup()
14     return None, None

```

Listing 2: Sandboxed Git Clone (src/tools/repo_{tools}.py)

Security Considerations:

- No execution of cloned code — analysis is static (AST) only.
- Git authentication errors are caught and reported, not exposed to shell.
- Timeouts prevent denial-of-service via malicious repos with large history.

2 Known Gaps and Concrete Plans

2.1 Judicial Layer: Current Status and Plan

Current Implementation:

- **Persona Prompts:** Distinct system prompts for Prosecutor (adversarial), Defense (optimistic), and TechLead (pragmatic) are defined in src/nodes/judges.py.
- **Structured Output:** Judges use .with_structured_output(JudicialOpinion) with retry logic (max 3 attempts) for malformed responses.
- **Parallel Execution:** Graph wiring supports fan-out to all three judges concurrently (pending full integration).
- **Evidence Formatting:** _format_evidence() converts collected forensic evidence into judge-consumable context blocks.

Known Gaps:

1. **Rubric Integration:** Judges currently use hardcoded criteria; need dynamic loading from week2_rubric.json via the Targeting Protocol.
2. **Persona Distinctiveness:** Prompts are distinct but not yet validated for >50% text overlap (risk of “Persona Collusion”).
3. **Error Resilience:** If one judge fails, the graph currently halts; need conditional edges to allow partial progress.
4. **Citation Enforcement:** Judges should cite specific evidence keys; currently optional in schema.

Concrete Plan (Completion by Final Submission):

1. **Week 1 (Remaining):**
 - Implement rubric loader in ContextBuilder node to inject judicial_logic per criterion.
 - Add prompt overlap checker (simple diff) to flag potential Persona Collusion during development.
2. **Week 2 (Final):**
 - Add conditional edges: if judge fails after retries, log warning and continue with available opinions.
 - Enforce cited_evidence as required field in JudicialOpinion; add validation in _invoke_judge().
 - Write integration tests: mock evidence → verify distinct judge outputs for same input.

2.2 Synthesis Engine: Current Status and Plan

Current Implementation:

- **Deterministic Logic:** `src/nodes/justice.py` implements score resolution via hardcoded Python (not LLM), with rules for variance handling.
- **Constitutional Overrides:** `_apply_overrides()` enforces `security_override` and `fact_supremacy` rules.
- **Report Generation:** `generate_markdown_report()` produces structured Markdown with Executive Summary, Criterion Breakdown, and Remediation Plan.
- **Fallback Handling:** If no opinions received, returns minimal report with explanatory message.

Known Gaps:

1. **Variance Re-evaluation:** `variance_re_evaluation` rule is stubbed but not fully implemented (needs evidence re-examination logic).
2. **Dissent Summaries:** Generated but not yet integrated with judge opinion citations for traceability.
3. **Remediation Specificity:** Current remediation is criterion-level; spec requires file-level instructions.
4. **LangSmith Integration:** Tracing is enabled but not yet verified end-to-end for synthesis node.

Concrete Plan (Completion by Final Submission):

1. Week 1 (Remaining):

- Implement `variance_re_evaluation`: re-query judges with highlighted evidence discrepancies when variance > 2.
- Enhance dissent summaries to include specific evidence keys cited by each judge.

2. Week 2 (Final):

- Extend `CriterionResult.remediation` to accept file-path + line-number suggestions from TechLead opinions.
- Add LangSmith trace verification: ensure synthesis node inputs/outputs are logged with criterion-level granularity.
- Write end-to-end test: mock full judge output → verify Markdown report structure and override application.

3 Planned StateGraph Flow

3.1 Architecture Overview

The Automaton Auditor implements a **two-level parallel fan-out/fan-in** topology using LangGraph's `StateGraph`:

1. **Level 1 (Detectives):** Three forensic agents run concurrently to collect evidence from GitHub repo and PDF report.
2. **Synchronization:** `EvidenceAggregator` node waits for all detectives before proceeding.
3. **Level 2 (Judges):** Three judicial personas evaluate the same evidence in parallel for each rubric criterion.
4. **Synthesis:** `ChiefJustice` node applies deterministic rules to resolve conflicts and generate final report.

3.2 StateGraph Diagram




Figure 1: StateGraph Topology: Two-level parallel execution with synchronization barriers. Detectives collect evidence concurrently; Judges evaluate the same evidence in parallel; Chief Justice synthesizes final verdict.

3.3 Key Execution Characteristics

- **Parallelism:** Detectives and Judges execute concurrently via LangGraph’s thread pool, reducing total latency.
- **State Reduction:** evidences: `Annotated[Dict, operator.ior]` and opinions: `Annotated[List, operator.add]` ensure thread-safe accumulation.
- **Error Handling:** Conditional edges (planned) will allow partial progress if individual nodes fail.
- **Observability:** LangSmith tracing captures each node’s inputs/outputs for debugging and auditability.

4 Conclusion and Next Steps

4.1 Summary of Progress

- **Foundation:** Pydantic state schemas, AST-based forensic tools, and sandboxed git operations are implemented and tested.
- **Graph Wiring:** Detective layer fan-out/fan-in is wired; judicial layer structure is defined.
- **Observability:** LangSmith integration enables end-to-end tracing of agent decisions.
- **Testing:** Unit tests for state models and tool functions; integration tests pending judicial layer completion.

4.2 Immediate Next Steps (This Week)

1. Complete rubric loading and Targeting Protocol integration in `ContextBuilder`.
2. Finalize judicial layer: dynamic criterion evaluation, persona distinctiveness validation, error-resilient graph edges.
3. Implement variance re-evaluation and file-level remediation in `ChiefJustice`.
4. Write end-to-end integration tests with mocked evidence and judge outputs.
5. Generate self-audit report and capture LangSmith trace for submission.

4.3 Risk Mitigation

- **LLM API Limits:** Fallback to smaller models (Ollama) if Groq/OpenAI rate limits are hit.
- **Complexity:** Modular node design allows incremental testing; each layer can be validated independently.
- **Time:** Prioritize core functionality (evidence collection + basic judging) before advanced features (vision, dissent summaries).

A Appendix: Representative Code Snippets

A.1 AgentState Definition with Reducers

```

1 class AgentState(TypedDict):
2     repo_url: str
3     pdf_path: str
4     rubric_dimensions: List[Dict]
5     # Parallel-safe reducers
6     evidences: Annotated[Dict[str, List[Evidence]], operator.ior]
7     opinions: Annotated[List[JudicialOpinion], operator.add]
8     final_report: Optional[AuditReport]
9     errors: Annotated[List[str], operator.add]

```

Listing 3: src/state.py: Parallel-safe state definition

A.2 AST Visitor for Graph Structure Detection

```

1 class ASTVisitor(ast.NodeVisitor):
2     def visit_ClassDef(self, node):
3         self.classes.append({
4             "name": node.name,
5             "bases": [ast.unparse(b) for b in node.bases],
6             "lineno": node.lineno
7         })
8         self.generic_visit(node)
9
10    def visit_Call(self, node):
11        try:
12            self.function_calls.append(ast.unparse(node.func))
13        except:
14            pass
15        self.generic_visit(node)

```

Listing 4: src/tools/repo_tools.py : ASTVisitor class

References

- [1] LangChain AI. *LangGraph Documentation*. 2026. <https://langchain-ai.github.io/langgraph/>
- [2] Samuel Colvin et al. *Pydantic: Data validation using Python type annotations*. 2026. <https://docs.pydantic.dev/>
- [3] Groq Inc. *Groq API Documentation*. 2026. <https://console.groq.com/docs>
- [4] Ollama. *Local Large Language Models*. 2026. <https://ollama.com>
- [5] LangChain AI. *LangSmith: Observability for LLM Applications*. 2026. <https://docs.smith.langchain.com/>